

# ASSIGNMENT 2

## Part 3 (Build and Train an RNN on MNIST)

After training our network, we had a Loss = 0.069390 and Training Accuracy = 0.98000. Also, the Testing Accuracy = 0.9804. Which is a very good result.

However, after we changed line 35 in our code to `rnn_cell.BasicRNNCell(nHidden)` in order to use LSTM and GRU instead of RNN. We had a Loss = 0.051129 and Training Accuracy = 0.99000 with Testing Accuracy = 0.9769. Also a very good outcome. However, the formal is better.

Comparing these two results to training using convnet in assignment 1. In assignment 1, We had a Training Accuracy = 0.98 with Testing Accuracy = 0.9883.

All three networks gives great outcome, however, training using convnet has the highest Testing Accuracy but it takes the highest c time.

```
In [2]: import tensorflow as tf
        from tensorflow.python.ops import rnn, rnn_cell
        import numpy as np

        from tensorflow.examples.tutorials.mnist import input_data

        mnist = input_data.read_data_sets("/tmp/data/", one_hot = True)

        learningRate = .0001
        trainingIters = 50
        batchSize = 100
        displayStep = 10

        nInput = 28 #we want the input to take the 28 pixels
        nSteps = 28 #every 28
        nHidden = 128 #number of neurons for the RNN
        nClasses = 10 #this is MNIST so you know

        x = tf.placeholder('float', [None, nSteps, nInput])
        y = tf.placeholder('float', [None, nClasses])

        weights = {
            'out': tf.Variable(tf.random_normal([nHidden, nClasses]))
        }

        biases = {
            'out': tf.Variable(tf.random_normal([nClasses]))
        }

        def RNN(x, weights, biases):
            x = tf.transpose(x, [1,0,2])
```

```

    x = tf.reshape(x, [-1, nInput])
    x = tf.split(x, nSteps, 0) #configuring so you can get it as needed for the 28 pixels

    lstmCell = rnn_cell.BasicLSTMCell(nHidden, state_is_tuple=True) #find which lstm to use in the documentation

    outputs, states = rnn.static_rnn(lstmCell, x, dtype=tf.float32) #for the rnn where to get the output and hidden state

    return tf.matmul(outputs[-1], weights['out']) + biases['out']

pred = RNN(x, weights, biases)

#optimization
#create the cost, optimization, evaluation, and accuracy
#for the cost softmax_cross_entropy_with_logits seems really good
cost = tf.reduce_mean( tf.nn.softmax_cross_entropy_with_logits(logits=pred, labels=y))
optimizer = tf.train.AdamOptimizer(learningRate).minimize(cost)

correctPred =tf.equal(tf.argmax(pred, 1), tf.argmax(y, 1))
accuracy = tf.reduce_mean(tf.cast(correctPred, 'float'))

init = tf.global_variables_initializer()

with tf.Session() as sess:
    sess.run(init)

    # while step* batchSize < trainingIters:
    for step in range(trainingIters):

        for _ in range(int(mnist.train.num_examples / batchSize)):
            batchX, batchY = mnist.train.next_batch(batchSize) #mnist has a way to get the next batch
            batchX = batchX.reshape((batchSize, nSteps, nInput))

            sess.run(optimizer, feed_dict={x: batchX, y:batchY})

        if step % displayStep == 0:
            acc = sess.run(accuracy, feed_dict={x: batchX, y:batchY})
            loss = sess.run(cost, feed_dict={x: batchX, y:batchY})
            print("Iter " + str(step) + ", Minibatch Loss= " +
                  "{:.6f}".format(loss) + ", Training Accuracy= " +
                  "{:.5f}".format(acc))

    print('Optimization finished')

    testData = mnist.test.images.reshape((-1, nSteps, nInput))
    testLabel = mnist.test.labels
    print("Testing Accuracy:", sess.run(accuracy, feed_dict={x: testData, y:testLabel}))

```

```

Extracting /tmp/data/train-images-idx3-ubyte.gz
Extracting /tmp/data/train-labels-idx1-ubyte.gz
Extracting /tmp/data/t10k-images-idx3-ubyte.gz
Extracting /tmp/data/t10k-labels-idx1-ubyte.gz

```

```

W1105 06:51:23.760756 3676 deprecation.py:323] From <ipython-input-2-cf8d
60855ef7>:35: BasicLSTMCell.__init__ (from tensorflow.python.ops.rnn_cell_
impl) is deprecated and will be removed in a future version.
Instructions for updating:
This class is equivalent as tf.keras.layers.LSTMCell, and will be replaced
by that in Tensorflow 2.0.
W1105 06:51:23.761754 3676 deprecation.py:323] From <ipython-input-2-cf8d
60855ef7>:37: static_rnn (from tensorflow.python.ops.rnn) is deprecated an
d will be removed in a future version.
Instructions for updating:
Please use `keras.layers.RNN(cell, unroll=True)`, which is equivalent to t
his API
W1105 06:51:23.799687 3676 deprecation.py:506] From C:\Users\oyeoy\Anacon
da3\lib\site-packages\tensorflow\python\ops\init_ops.py:1251: calling Vari
anceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is d
eprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to
the constructor
W1105 06:51:23.815614 3676 deprecation.py:506] From C:\Users\oyeoy\Anacon
da3\lib\site-packages\tensorflow\python\ops\rnn_cell_impl.py:738: calling
Zeros.__init__ (from tensorflow.python.ops.init_ops) with dtype is depreca
ted and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing it to
the constructor
W1105 06:51:24.466904 3676 deprecation.py:323] From <ipython-input-2-cf8d
60855ef7>:46: softmax_cross_entropy_with_logits (from tensorflow.python.op
s.nn_ops) is deprecated and will be removed in a future version.
Instructions for updating:

Future major versions of TensorFlow will allow gradients to flow
into the labels input on backprop by default.

See `tf.nn.softmax_cross_entropy_with_logits_v2`.

```

```

Iter 0, Minibatch Loss= 0.490804, Training Accuracy= 0.86000
Iter 10, Minibatch Loss= 0.046969, Training Accuracy= 0.98000
Iter 20, Minibatch Loss= 0.043729, Training Accuracy= 0.97000
Iter 30, Minibatch Loss= 0.063595, Training Accuracy= 0.98000
Iter 40, Minibatch Loss= 0.069390, Training Accuracy= 0.98000
Optimization finished
Testing Accuracy: 0.9804

```

In [ ]:

In [ ]: